

PROBLEM

AI agents are becoming operational workers, but real work is not a chat transcript.

When an agent touches production systems, teams need to know what happened, what evidence was used, what risk was found, who approved it, and how work can be resumed safely.

AgentOps Case Control Room

Turn AI-agent work into a UiPath-governed case with evidence, risk, approval, robot work items, and human handoff.
Business impact: safer AI operations for release, support, IT, and automation teams.

UiPath AgentHack · AgentOps Case Control Room · Demo Artifact

Turn AI-agent work into a governed UiPath case.

This prototype shows how a human, an AI coding agent, UiPath robots, APIs, and an approval owner move through one auditable operations case instead of leaving work trapped in a chat transcript.

CASE

CASE-AI-OPS-001 · AI Agent Operations Exception Review

UIPATH TARGET

Maestro Case

CURRENT STATUS

release rejected until evidence fixed

MAX RISK

critical

11

evidence events in the UiPath case packet

3

human approval gates before sensitive action

3

robot/API evidence collection steps

1

production action blocked by policy

Maestro Case Flow

?

What did the agent do?

!

Was a risky action blocked?

✓

Who approved the final decision?

evt-0007

affects payments and has failing verification.

Attempted production deployment was blocked by policy before execution.

evt-0008

policy before execution.

Human requested more evidence instead of approving the production deployment.

evt-0010

Service owner rejected the release until the retry regression is fixed.

evt-0011

Generated a case handoff report with cited evidence IDs and final rejection decision.

Human Approval Queue

EVENT	PHASE	RISK	DECISION	REASON
evt-0007	execution	CRITICAL	blocked_by_policy	Attempted production deployment was blocked by policy before execution.
evt-0008	approval	MEDIUM	needs_more_evidence	Human requested more evidence instead of approving the production deployment.
evt-0010	approval	LOW	rejected	Service owner rejected the release until the retry regression is fixed.

Final Handoff


The release is rejected until the retry regression is fixed. The case remains useful after the AI conversation ends because every decision points back to a stable evidence event ID.

SOLUTION

A UiPath-governed case for human-AI operations.

- Human opens an exception case.
- AI agent creates a read-only investigation plan.
- Robot gathers ticket, pull request, and owner evidence.
- Policy gateway blocks risky production action.
- Action Center-style tasks route approval to humans.

Core thesis: agents can move fast, but UiPath makes the work governable.



ARCHITECTURE

Maestro Case

Case stages: intake, agent investigation, robot/API evidence, risk review, approval, handoff.

Robot Work Items

Evidence collection is modeled as queue-ready robot work items.

Action Center

Human approval tasks preserve evidence, risk, and decision context.

Implementation artifacts include a BPMN-style process blueprint, case data model, Action Center form schema, robot queue model, and simulated case run.

uipath-package/maestro-process.bpmn

runtime/robot-work-items.json

action-center/action-center-tasks.json

Claim boundary: verified local prototype and UiPath-ready implementation package.

Repeatable local verification is included.

- 26 shared AgentOps events
- 3 cases
- 5 approval gates
- 1 blocked production action
- 2 robot work items
- 3 Action Center tasks

```
cd uipath-agenthack
```

```
bash scripts/run_uipath_local_checks.sh
```

```
verify_ok
```

```
uipath_verify_ok
```

WHY NOW

Human-AI Operations Control Plane

One event stream for humans, AI agents, robots, APIs, approvals, risks, costs, and evidence. This is the shared prototype behind the UiPath, Splunk, FIND EVIL, Google Cloud, and Microsoft hackathon lanes.

Events 26	Cases 3	Approvals 5	Blocked 1	High/Critical 3	Est. Cost \$4.86
--------------	------------	----------------	--------------	--------------------	---------------------

UiPath Case Workbench

1

Intake
Human opens the case
evt-0001

none

2

Agent Investigation
AI agent plans and tests read-only
evt-0002, evt-0005

medium

3

Robot Evidence
UiPath robot and APIs gather records
evt-0003, evt-0004, evt-0009

low

4

Risk Review
Control plane escalates and blocks risk
evt-0006, evt-0007

critical

5

Human Approval
Humans ask for evidence and reject release
evt-0008, evt-0010

medium

6

Handoff
Report is generated from evidence IDs
evt-0011

none

Human Approval Queue

evt-0007 **codex-ops-agent**
Attempted production deployment was blocked by policy before execution.
BLOCKED BY POLICY

evt-0008 **operations-lead**
Human requested more evidence instead of approving the production deployment.
NEEDS MORE EVIDENCE

evt-0010 **service-owner**
Service owner rejected the release until the retry regression is fixed.
REJECTED

evt-0018 **security-analyst**
Human approved containment of the reported message and requested no endpoint isolation.
APPROVED

evt-0025 **support-manager**
Human approved the workflow with a rule that expensive model calls are escalation-only.
APPROVED

Contest Adapters

UiPath case_packet.json	Splunk hec_events.jsonl	FIND EVIL dfir_report.md	Google gemini_mcp_workflow.json	Microsoft copilot_architecture.md
----------------------------	----------------------------	-----------------------------	------------------------------------	--------------------------------------

CASE-AI-OPS-001

critical

AI Agent Operations Exception Review

11 events

3 approvals

1 blocked

\$0.006

1. evt-0005 Regression test failed around duplicate retry suppression.

2. evt-0006 Production payment behavior change plus failing regression test.

3. evt-0007 Production deployment is destructive and requires human approval.

4. evt-0008 Human requested more evidence instead of approving the production deployment.

CASE-DFIR-002

high

Evidence-Locked Phishing Investigation

8 events

1 approvals

0 blocked

\$0.012

1. evt-0015 Initial model answer overclaimed malware execution before process evidence was present.

2. evt-0016 Unsupported certainty detected in analyst draft.

3. evt-0017 Hash lookup returned low-confidence phishing kit association; raw indicator was redacted in report output.

CASE-CLOUD-003

medium

Human-Controlled Cloud Agent Workflow

7 events

1 approvals

0 blocked

\$4.842

1. evt-0024 Projected monthly model spend exceeded the manager's prototype budget.

Event Timeline							Case	All cases	
EVENT	PHASE	ACTOR	TYPE	STATUS	RISK	SUMMARY			
evt-0001	intake	human operations-lead	task_start	info	none	Human opened an AI-agent operations case after a release bot proposed a production change.			
evt-0002	planning	ai_agent codex-ops-agent	plan_update	success	low	Agent created a read-only investigation plan with explicit approval gates before any production action.			
evt-0003	evidence_collection	robot uipath-evidence-robot	robot_task	success	low	Robot gathered the change ticket, linked pull request, deployment window, and service owner.			
evt-0004	evidence_collection	api github-api	api_call	success	low	API confirmed the pull request touches payment retry behavior and has one approving review.			
evt-	investigation	ai_agent	tool_call	failed	medium	Agent found a failing retry regression test that was not mentioned in the change			

AgentOps Case Control Room shows how human, AI, robot, and API work can be reviewed as one case. That is the bridge between a

prototype on a laptop and agentic automation that enterprises can trust.

Next: import the generated process, queue, and approval models into UiPath Automation Cloud / Maestro and verify a live run.

0005		codex-ops-agent		failure	medium	ticket.
evt-0006	risk_review	ai_agent codex-ops-agent	risk_signal	warning	high	Case risk escalated because the proposed release affects payments and has failing verification.
evt-0007	execution	ai_agent codex-ops-agent	tool_call	blocked	critical	Attempted production deployment was blocked by policy before execution.
evt-0008	approval	human operations-lead	approval_gate	success	medium	Human requested more evidence instead of approving the production deployment.
evt-0009	evidence_collection	robot ulpath-evidence-robot	robot_task	success	low	Robot requested service-owner signoff and attached the failed test evidence.
evt-0010	approval	human service-owner	approval_gate	success	low	Service owner rejected the release until the retry regression is fixed.
evt-0011	handoff	system agentops-recorder	handoff_report	success	none	Generated a case handoff report with cited evidence IDs and final rejection decision.
evt-0012	intake	human security-analyst	task_start	info	none	Security analyst opened a suspicious email investigation with attached endpoint telemetry.
evt-0013	planning	ai_agent dfir-analysis-agent	plan_update	success	low	Agent proposed an evidence-first DFIR plan: preserve artifacts, list hypotheses, then verify each claim.
evt-0014	evidence_collection	robot sift-collector	evidence_captured	success	low	Collected email headers, attachment hash, browser download timeline, and endpoint process list.
evt-0015	investigation	ai_agent dfir-analysis-agent	ai_call	warning	medium	Agent generated an initial hypothesis but flagged one claim as unsupported by current evidence.
evt-0016	risk_review	ai_agent dfir-analysis-agent	risk_signal	warning	high	System marked the draft finding as provisional because execution evidence was missing.
evt-0017	evidence_collection	api threat-intel-api	api_call	success	medium	Hash lookup returned low-confidence phishing kit association; raw indicator was redacted in report output.
evt-0018	approval	human security-analyst	approval_gate	success	low	Human approved containment of the reported message and requested no endpoint isolation.
evt-0019	handoff	system agentops-recorder	handoff_report	success	none	Generated DFIR report with evidence IDs, unsupported claims, and human containment decision.
evt-0020	intake	human support-manager	task_start	info	none	Support manager requested an agent that answers customer questions and escalates uncertain cases.
evt-0021	planning	ai_agent gemini-ops-navigator	plan_update	success	low	Agent created a workflow plan using retrieval, MCP tools, cost guardrails, and human escalation.
evt-0022	evidence_collection	api knowledge-base-mcp	api_call	success	low	MCP retrieval returned the current refund policy and source URL for citation.
evt-0023	execution	ai_agent gemini-ops-navigator	ai_call	success	low	Agent drafted a customer answer grounded in the retrieved refund policy.
evt-0024	risk_review	ai_agent gemini-ops-navigator	cost_signal	warning	medium	Cost guardrail warned that the high-quality model should be reserved for escalations.
evt-0025	approval	human support-manager	approval_gate	success	low	Human approved the workflow with a rule that expensive model calls are escalation-only.
evt-0026	handoff	system agentops-recorder	task_end	success	none	Closed cloud-agent workflow case with source citations, cost guardrail, and human approval captured.